

offline AI.

Qwen 3.5 lets you run a real AI model on your phone. No internet. No subscription. No data leaving your device. Download once, use forever.

What Is Qwen 3.5?

- Alibaba's Small Model Series, built for on-device use
- 4 sizes: 0.8B (ultra light), 2B (fast), 4B (sweet spot), 9B (near cloud quality)
- Open source (Apache 2.0), free on Hugging Face

Cloud AI vs Offline AI

BEFORE

- ✗ Needs internet always
- ✗ Data leaves your device
- ✗ Monthly subscription
- ✗ Latency on every request
- ✗ Unusable offline



AFTER

- ✓ Download once, run forever
- ✓ Data stays on device
- ✓ Completely free
- ✓ Instant local response
- ✓ Works in airplane mode

Not Just **Text** Chat

- 01** Text Generation: Full conversations and content creation
- 02** Vision AI: Analyze images directly on device
- 03** Voice Transcription: Whisper-powered speech to text
- 04** Document Analysis: Read and understand PDFs and docs

How To Run It

IPHONE


- App: Off Grid / PocketPal AI
- Needs: iPhone 12+, iOS 17+
- Best: iPhone 15 Pro (8GB)
- Speed: 20-30 tokens/sec
- Tip: Set KV cache to q4_0



ANDROID

- App: Off Grid (Play Store)
- Needs: 6GB RAM, ARM64
- Best: Snapdragon 8 Gen 2+
- Speed: 10-30 tokens/sec
- Works on \$300 phones

Pick Your Model




Real Use Cases

- Developer debugging code on a flight with no Wi-Fi
- Student in a rural area learning from an AI tutor offline
- Field worker getting instant answers without cell coverage
- Privacy-first user who never wants data leaving their device
- Apps that work without internet dependency

Setup: iPhone


- App Store: 'Off Grid' or 'PocketPal AI'
- Both free, no account needed
- Tap model browser
- Models filtered by your device RAM
- Start with Qwen 3 0.6B or 1.5B



```
App Store > Search 'Off Grid'  
✓ Installed (free)  
  
> Open Off Grid > Model Browser  
Qwen 3 0.6B ..... 80MB  
Qwen 3 1.5B ..... 640MB  
Qwen 3 4B ..... 2.1GB  
  
> Download over Wi-Fi  
✓ Done. Start chatting instantly.
```

Setup: Android

- Open Play Store, search 'Off Grid'
- Install (free, open source, MIT)
- Browse recommended models
- Pick based on your RAM (6GB+)
- Turn on airplane mode to verify



```
Play Store > Search 'Off Grid'  
✓ Installed (open source, MIT)  
> Browse recommended models  
RAM: 6GB+ needed  
Download: 80MB to 4GB  
> Airplane mode: ON  
✓ Still works. Zero internet.
```

Pro Tip: Speed Boost

- After your first chat, go to Settings
- Find 'KV Cache Quantization'
- Switch from f16 to q4_0
- This roughly triples inference speed with minimal quality loss
- First message is always slower (model loads into memory)
- Subsequent messages are much faster



**AI that works without the
internet is not a feature. It's the
future.**

Follow for More!



Shaïd

GenAI Architect & Builder

AI Products & Automation | Ship AI. Share Insights.

#GenerativeAI #EdgeAI #OnDeviceAI #Qwen #OfflineAI