

Opus 4.7
is here.

Anthropic just replaced **Opus 4.6.
Devs had been yelling about
regressions for weeks. Here's what
4.6 got wrong — and what **4.7**
actually fixes.**

The Announcement

April 16, 2026 — Anthropic confirms the upgrade

Introducing Claude Opus 4.7, our most capable Opus model yet.

It handles long-running tasks with more rigor, follows instructions more precisely, and verifies its own outputs before reporting back.

Why 4.6 Was Being Replaced

- Fabricated data: prices, file sizes, API capabilities
- Panic loops: 3-call tasks turning into 30+ tool calls
- Research subagents returned fake data, main agent trusted it
- BridgeBench hallucination test: 83% dropped to 68%
- Reasoning depth fell ~67% by late February
- Broke GUI rules, took JavaScript shortcuts instead

Real Production Failures

- User burned 80% of Max 20x weekly budget in 2.5 days
- CLAUDE.md rules loaded every prompt, ignored anyway
- Flip-flopped: "keep test" then "revert" then "@Disabled"
- 8 spreadsheet rows corrupted in a translation QA session
- Ran ONE targeted test, missed 3 regression categories
- Forgot MCP servers mid-session, re-researched what it knew



**Claude has regressed to the point it cannot be trusted to perform complex engineering.
— AMD Senior Director, on 4.6**

What 4.7 Actually Fixes

BEFORE

- ✗ Gave up on hard multi-step tasks
- ✗ Missed its own logical faults
- ✗ Tool errors piled up in agents
- ✗ Low-res vision on diagrams
- ✗ Ambiguous instruction following



AFTER

- ✓ Pushes through complex problems
- ✓ Verifies outputs before replying
- ✓ 1/3 fewer tool errors in agents
- ✓ 3x higher image resolution
- ✓ Precise, deliberate execution

The Benchmark Receipt

Opus 4.7 beats 4.6, GPT-5.4, and Gemini 3.1 Pro on most rows


You can hand off your hardest work with less supervision.

Opus 4.7

	Opus 4.7	Opus 4.6	GPT-5.4	Gemini 3.1 Pro	Mythos Preview
Agentic coding SWE bench Pro	64.3%	53.4%	57.7%	54.2%	77.8%
Agentic coding SWE bench Verified	87.6%	80.8%	—	80.6%	93.9%
Agentic terminal coding Terminal Bench 2.0	69.4%	65.4%	75.1% self-reported harness	66.5%	82.0%
Multidisciplinary reasoning Humanity's Last Exam	46.9% no tools	40.0% no tools	42.7% no tools (Pro)	44.4% no tools	56.8% no tools
	54.7% with tools	53.3% with tools	58.7% with tools (Pro)	51.4% with tools	64.7% with tools
Agentic search BrowseComp	79.3%	83.7%	89.3% Pro	85.9%	86.9%
Scaled tool use MCP Atlas	77.3%	75.8%	68.1%	73.9%	—
Agentic computer use OSWorld Verified	78.0%	72.7%	75.0%	—	79.6%
Agentic financial analysis Finance Agent v1	64.4%	60.1%	61.5% Pro	59.7%	—
Cybersecurity vulnerability reproduction CyberOps	73.1%	73.8%	66.3%	—	83.1%
Graduate-level reasoning GPQA Diamond	94.2%	91.3%	94.4% Pro	94.3%	94.6%
Visual reasoning Char-Kis Reasoning	82.1% no tools	69.1% no tools	—	—	86.1% no tools
	91.0%	84.7%	—	—	93.2%

The Coding Jump

Double-digit gains on SWE-bench Pro, CursorBench, and Visual Acuity



Vision: The Biggest Leap

BEFORE

- ✗ Visual acuity: 54.5% on XBOW
- ✗ Struggled with chemical diagrams
- ✗ Computer-use automation failed often
- ✗ 21% more errors on document reasoning



AFTER

- ✓ Visual acuity: 98.5% on XBOW
- ✓ Reads circuit and chemical diagrams
- ✓ Reliable computer-use automation
- ✓ 2,576 px edge (~3.75 MP) image input

New: **/ultrareview** in Claude Code

- Deep-dive code review: architecture, security, performance, maintainability
- Pro and Max: 3 complimentary sessions
- Built for multi-file refactors and race-condition debugging
- Uses 4.7's upgraded long-context retrieval (still 1M tokens)
- Multi-needle retrieval now holds deep inside documents

What **Reviewers** Noticed

- 01** +13% on 93-task coding benchmark
- 02** +14% on multi-step workflows at fewer tokens
- 03** 77.1% on ARC-AGI-2 — beats Gemini 3 Pro
- 04** 90.9% on BigLaw Bench legal analysis
- 05** 94.2% on GPQA Diamond graduate reasoning

The Honest Catch

- Anthropic admits 4.7 still trails unreleased Mythos
- Tighter instruction following may break old prompts
- It's incremental, not a paradigm shift
- But same \$5 / \$25 per M tokens as 4.6
- If you're on 4.6 today, this is free lift

The Takeaway

- 4.6 broke user trust. 4.7 is Anthropic's apology in code
- Coding, vision, agent reliability all sharper
- Same price, stronger model — easy swap
- Retune prompts if instruction following gets too strict
- Ship the upgrade. Keep Mythos on your radar.

Who Wins with 4.7

- Engineers: multi-file refactors, race-condition debugging
- Legal teams: Harvey.ai live — drafting, deal management at 90.9%
- Finance analysts: SOTA on GDPval-AA economic work
- Agent builders: 1/3 fewer tool errors, reliable computer use
- Researchers: reads circuits, UML, chemical structures
- The Neuron's live warts-and-all test verdict: worth the switch

Follow for More!



Shaïd

GenAI Architect & Builder

AI Products & Automation | Ship AI. Share Insights.

#ClaudeOpus #Anthropic #AI #GenAI #AICoding