


\$4,515
in tokens.

I built a **local dashboard that reads my Claude Code session files and turns them into spend, prompt, and project analytics. The numbers stopped me cold.**

My Actual Dashboard

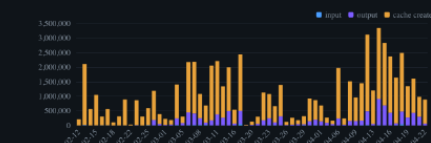
241 sessions, 22,644 turns, 1.9B cache reads, \$4,515 in API-equivalent cost.



What do these numbers mean?click to expand

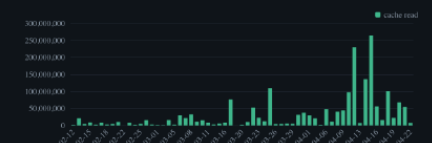
Your daily work

Tokens you paid for: what you sent (input), what Claude wrote (output), and what got stored for re-use (cache create).



Daily cache reads

Cache reads are cheap re-uses of things Claude already saw (like your CLAUDE.md). They cost ~10x less than regular input tokens — high numbers here are a good thing.

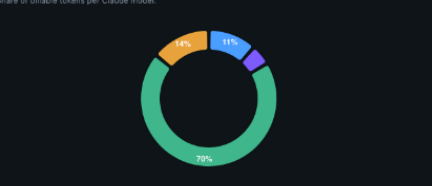


Tokens by project

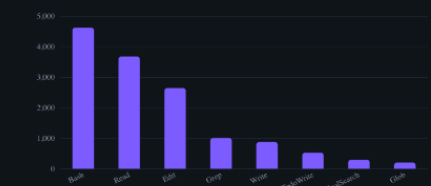


Token usage by model

Share of billable tokens per Claude model.



Top tools (by call count)



Recent sessions

all →

STARTED	PROJECT	tokens
2024-04-21 05:124	Nisha	42.4K
2024-04-14 15:156	Nisha	202.1K
2024-04-21 14:119	Nisha	50.7K
2024-04-22 07:409	Nisha	3K
2024-04-21 05:117	Nisha	124K
2024-04-21 08:446	Nisha	16.9K
2024-04-21 07:425	Nisha	67.4K
2024-04-21 07:441	Nisha	1.9K
2024-04-21 05:126	Nisha	0
2024-04-21 04:143	Nisha	19.7K

The numbers that **stopped me**

- 01** 241 sessions all-time
- 02** 22,644 conversation turns
- 03** 9.9 MILLION output tokens
- 04** 1.9 BILLION cache-read tokens
- 05** \$4,515 API-equivalent cost



**1.9 billion cache reads vs 332K
fresh input. That's why Claude
Code feels affordable — you're
paying for a sliver of what
you're actually reading.**

Before vs **After** the dashboard

BEFORE

- ✗ No idea what I was spending
- ✗ No idea which prompts were expensive
- ✗ Trusted the monthly bill blindly



AFTER

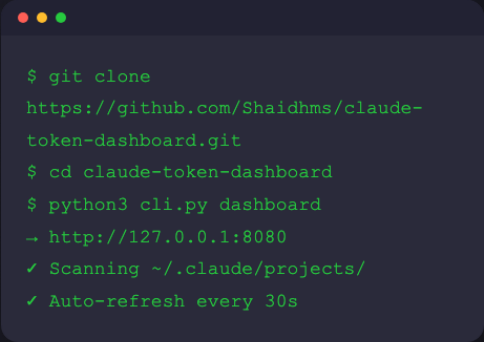
- ✓ Exact cost per session
- ✓ Most expensive prompts ranked
- ✓ Tips to reduce token usage

7 tabs. One full picture.

- Overview — totals, daily charts, cost
- Prompts — your most expensive ones
- Sessions — turn-by-turn breakdown
- Projects — where tokens go
- Skills — which skills you use most
- Tips — how to reduce usage
- Settings — API / Pro / Max pricing

Set up in 30 seconds

- Pure Python stdlib
- No pip install
- No Node.js
- No build step
- 100% local. Nothing leaves your machine.



```
$ git clone
https://github.com/Shaidhms/claude-
token-dashboard.git
$ cd claude-token-dashboard
$ python3 cli.py dashboard
→ http://127.0.0.1:8080
✓ Scanning ~/.claude/projects/
✓ Auto-refresh every 30s
```

What I actually learned

- Cache is the hero — 10x savings happen in the background
- Big tool results are the real cost driver
- One bad prompt can burn more than a whole week of small ones
- My Nisha project alone: 1.74B cache reads



**GitHub link in the description.
Clone it. See what you're
actually spending — and get
tips to reduce it.**

Follow for more!



Shaïd

GenAI Architect & Builder

Ship AI. Share Insights. github.com/Shaidhms/claude-token-dashboard

#ClaudeCode #AI #GenAI #Anthropic #BuildInPublic