


99%
Cheaper
with
Ollama.

You don't always need Opus. Most Claude Code tasks can run on **free local models — and two methods make it happen without breaking your workflow.**

2 Methods to Run Free




Method 1:

Ollama Local

Install Ollama, pull a model, point Claude Code at local endpoint.


Zero API cost. Runs offline.



```
$ ollama pull llama3.2
✓ Model ready at localhost:11434
$
ANTHROPIC_BASE_URL=http://localhost:11434
✓ Claude Code now using local model
```


Run Models **Locally Free**

Ollama.com — download and run open-source LLMs on your own machine



Method 2: Open Router

Free cloud models via Open Router.

No local GPU needed.

Works on any machine.

```
$ # Set Open Router as base URL
ANTHROPIC_BASE_URL=https://openrouter.ai/api
ANTHROPIC_API_KEY=<your-free-key>
✓ Free cloud model active
```

Open Source vs Opus

BEFORE

- ✗ Boilerplate and scaffolding code
- ✗ Quick one-file scripts
- ✗ Exploration and prototyping
- ✗ Repetitive refactoring tasks



AFTER

- ✓ Complex multi-file reasoning
- ✓ Architecture decisions
- ✓ Production-grade code review
- ✓ Debugging subtle logic errors

Where Open Source Wins

- 01** Generating boilerplate — no thinking required
- 02** Writing repetitive test files
- 03** Simple CLI scripts under 100 lines
- 04** Reformatting and data transformation
- 05** Exploration tasks you'll throw away
- 06** Anything with a clear template

Where Opus Still Wins

- Multi-file context with interdependencies
- Complex reasoning chains
- Catching subtle bugs in production code
- Architecture decisions with tradeoffs
- Anything where wrong = expensive



**Route the right task to the right
model. Use Ollama for volume,
Opus for value.**

Follow for More!



Shaïd

GenAI Architect & Builder

AI Products & Automation | Ship AI. Share Insights.

#Ollama #ClaudeCode #AI #OpenSource #GenAI #LocalAI #AITools